

# IRFAN SHARIF

irfansharif.io/blog — irfan@irfansharif.io — in/irfansharifm — github.com/irfansharif

## SUMMARY

---

- Experience with distributed systems, ML, storage and search engines, file systems, SAT solvers

## EDUCATION

---

**University of Waterloo** 2014 – 2019

*Bachelor of Applied Science in Honours Computer Engineering (with Distinction)*

## EXPERIENCE

---

**Cockroach Labs** *Toronto, ON*  
*Senior Software Engineer (Transactional KV)* 2019 – Present

- Leading admission control: overload protection, work prioritization, and inter-tenant fairness
- Built elastic CPU limiter using patched Go, protecting latencies from in-process background work
- Improved schema count scalability from  $O(10k)$  to  $O(1M)$ , re-architecting for multi-tenancy
- Worked on flow control for quorum writes, upgrades infrastructure, bazel-ifying monorepo, tracing

**LinkedIn** *Sunnyvale, CA*  
*Software Engineering Intern (Machine Learning Infrastructure)* 2018

- Worked on model serving, versioning, runtime and deployment across thousands of hosts

**Amazon** *Palo Alto, CA*  
*Software Engineering Intern (Search Relevance)* 2018

- Used sequence-to-sequence learning to generate alternative search suggestions, authored query reformulation modules to improve system recall, and incorporated OCR text into ranking models

**Cockroach Labs** *New York, NY*  
*Backend Engineering Intern (Storage, Distributed SQL)* 2016 – 2017

- Designed dedicated storage engine for Raft, contributed to TLA+ verified PreVote extension
- Implemented aggregations, sorting,  $n$ -way parallel joins, and row de-duplications for query engine

**Shopify** *Ottawa, ON*  
*Production Engineering Intern (Infrastructure)* 2016

- Built an auto-scaling build system with dynamic workload distribution handling 1000+ builds/day

## PUBLICATIONS

---

- [1] Rebecca Taft, **Irfan Sharif**, et al. CockroachDB: The Resilient Geo-Distributed SQL Database. In *Proceedings of the ACM SIGMOD International Conference on Management of Data*, 2020.
- [2] Nathan VanBenschoten, Arul Ajmani, Marcus Gartner, Andrei Matei, Aayush Shah, **Irfan Sharif**, et al. Enabling the Next Generation of Multi-Region Applications with CockroachDB. In *Proceedings of the ACM SIGMOD International Conference on Management of Data*, 2022.
- [3] **Irfan Sharif**. Hitchhiker's Guide to Transactions. In *Hydra: Computing Conference*, 2021.

## RESEARCH

---

**Waterloo Artificial Intelligence Institute** (*SAT/SMT solvers*) 2016

- Studied search space pruning, clustering and parallelization techniques under Prof. Vijay Ganesh

## PROJECTS

---

- **Kura** – An E2EE, synchronized, distributed file system with live-sharing and file streaming
- **Solver** – SAT solver library in Go (wraps around Google's Operational Research Tools)
- **CFilter** – Cuckoo Filter implementation with 770+ stars and 35+ forks on GitHub